

Promptist Reproduction Project Report

1. Executive Summary

This project successfully reproduced the results of Microsoft's **Promptist** (a reinforcement learning-based prompt optimization model) on a local Windows environment with limited VRAM (NVIDIA RTX 4050, 6GB).

We established a fully offline, quantization-enabled pipeline that allows for: 1. **Text Optimization**: Transforming simple prompts into aesthetic-heavy prompts using `microsoft/Promptist`. 2. **Image Generation**: Generating comparison images using `runwayml/stable-diffusion-v1-5`. 3. **VRAM Management**: Running both models sequentially on a 6GB card without OOM errors.

2. Environment Setup

The original repository required outdated dependencies incompatible with modern hardware (RTX 40 series) and Windows. We modernized the stack:

2.1 Core Dependencies

- **Python**: 3.11.0
- **PyTorch**: 2.6.0+cu124 (Upgraded from 1.x for CUDA 12 support)
- **BitsAndBytes**: 0.49.1 (Enabled 8-bit quantization on Windows)
- **Transformers**: Latest version compatible with PyTorch 2.x
- **Diffusers**: Patched local version to remove deprecated `huggingface_hub` imports ("cached_download" error).

2.2 Workarounds

- **TRXL**: The Reinforcement Learning library `trlx` has a hard dependency on `DeepSpeed`, which is difficult to build on Windows. We patched `setup.cfg` to remove this dependency, allowing the inference modules to load successfully.
-

3. Model Management

To ensure stability and offline capability, we bypassed runtime downloads and stored weights locally:

Model	Source	Local Path	Purpose
Promptist	microsoft/Promptist	./checkpoints/promptist	Text Optimization
CLIP	openai/clip-vit-la	./checkpoints/clip	Metric Evaluation (Optional)
Stable Diffusion	runwayml/stable-diffusion-hub	(Cached in HUB)	Image Generation

4. Implementation & Scripts

We developed a suite of scripts to verify, optimize, and generate results.

4.1 `reproduce_results_v2.py` (Inference Verification)

Purpose: Validates that the Promptist model loads and generates text correctly. - **Key Feature:** Uses BitsAndBytesConfig for 8-bit loading. - **Fixes:** Implements explicit attention_mask and pad_token_id handling to prevent GPT-2 generation errors. - **Output:** Prints the optimized version of "A futuristic city with flying cars".

4.2 `generate_report_images.py` (End-to-End Pipeline)

Purpose: Generates a side-by-side verification image while managing memory. - **Phase 1:** Loads Promptist (8-bit), optimizes prompt, saves text. - **Phase 2: Crucial VRAM Cleanup.** Explicitly deletes the LLM and runs `gc.collect()`/`cuda.empty_cache()` to free 6GB VRAM. - **Phase 3:** Loads Stable Diffusion (float16) to generate `baseline.png` and `optimized.png`.

4.3 `batch_test.py` & `batch_test_v2.py`

Purpose: Batch processing for evaluation on 7 diverse prompts. - `batch_test.py`: The initial script that performed both optimization and generation. - `save_batch_prompts.py`: Extracted the optimized prompt pairs to `batch_results/prompts.json` for analysis. - `batch_test_v2.py`: A specialized generation script that reads from the JSON file to reliably regenerate images without re-loading the LLM.

5. Workflow Diagram

```
graph TD A[Start] --> B(Setup Env & Patch TRLX); B --> C(Download Models Checkpoints); C --> D{Verification}; D -->|Step 1| E[reproduce_results_v2.py]; E -->|Success| F{Generation}; subgraph Single Image Pipeline F -->|Step 2| G[generate_report_images.py]; G --> H[Optimize Text]; H --> I[Clear VRAM]; I --> J[Generate Images]; end subgraph Batch Pipeline F -->|Step 3| K[batch_test.py / save_batch_prompts.py]; K --> L[Extract Prompts to JSON]; L --> M[batch_test_v2.py]; M --> N[Generate 14 Batch Images]; end
```

6. Results

6.1 Single Verification

Prompt: "A futuristic city with flying cars"

Baseline	Promptist Optimized
	
<i>Generic generation</i>	<i>Enhanced detail & style</i>

6.2 Batch optimization

Promptist effectively appends aesthetic modifiers. Example from our batch test:

Original: "A cat in the rain" **Optimized:** "cat in the rain, heavy rain, tears in rain, cinematic lighting, sharp focus, intricate, 8 k, detailed, art by artgerm and greg rutkowski and alphonse mucha"

All 14 batch images are stored in the `batch_results_v2/` directory.

7. Conclusion

The Promptist replication is complete. We have successfully: - [x] Overcome Windows/Hardware compatibility issues. - [x] Implemented a memory-efficient pipeline for 6GB VRAM. - [x] Verified results through both single-instance and batch testing.